

cogs9_proj

December 18, 2020

1 Cogs9 Project: For the SDPD in 2019, is there a significant difference in the likelihood of someone being stopped according to their race?

Group Name: CV GANG
By Andrew Cheng

1.1 Analysis

For our analysis method, we'll be using python pandas - link to repository-
<https://github.com/AndrewCheng2002/Cogs-9-Project>

```
[85]: #Modules
import numpy as np
import pandas as pd
import geopandas as gpd
import matplotlib.pyplot as plt
import math
import datetime

%matplotlib inline
plt.style.use('fivethirtyeight')
```

```
↳ -----
```

```
↳ last) ModuleNotFoundError Traceback (most recent call↳
```

```
<ipython-input-85-f6c182b9deb5> in <module>
      3 import pandas as pd
      4 import geopandas as gpd
----> 5 import matplotlib as matplotlib
      6 import math
```

```
7 import matplotlib.pyplot as plt
```

ModuleNotFoundError: No module named 'geoplot'

Data Collection: How do we get data to perform analysis on?

When answering a data science question, our first task is to gather the data itself. However, not just any data will suffice and in order to get an accurate, unbiased and ethical model, we need ensure that our data is sufficient and representative of the population. (San Diego in 2019) We needed to find a source that was reliable to ensure fairness and correct representation as well as one that was large enough to not be affected by outliers. After spending lots of time and deliberation we settled on using RIPA (An act passed by the government that requires police departments to publish their data regarding stops, arrests, etc.) to get data on [race](#) and [stop](#) information. We simply downloaded the data off the RIPA website which was conveniently in csv format. In addition to race and stop data we also need data of specific race demographics. In this case, we decided to use the US [census](#) which has information about the race distribution of the San Diego County in 2019. Even though we are analyzing the SDPD in 2019, which is not a census year (every decade), we felt that the predictive algorithms used by the US government were trustworthy and sufficient enough to get an accurate representation of the population demographic. Because there was no quick download function and the fact that we're only analyzing 7 races, we decided it was easier to just web-scrape the information and put it in a table.

```
[2]: #We made pandas read the RIPA csv files we downloaded
data_race_raw = pd.read_csv('ripa_race_data_sd.csv')
data_stops_raw = pd.read_csv('ripa_stops_data_sd.csv', low_memory = False)

#For the census we web-scraped the information and manually inserted in a table
data_census_race = pd.DataFrame({'percentage of population': [45,5.5,34.1,.9,12.
↪6,1.3,.6]},
index = ['White', 'Black/African American', 'Hispanic/Latino/a', 'Middle Eastern_
↪or South Asian', 'Asian', 'Native American', 'Pacific Islander'])
```

```
[3]: data_race_raw
```

```
[3]:
```

	stop_id	pid	race
0	2443	1	White
1	2444	1	White
2	2447	1	Hispanic/Latino/a
3	2447	2	Hispanic/Latino/a
4	2448	1	White
...
394970	356019	1	Black/African American
394971	356025	1	Black/African American
394972	356080	1	White
394973	356300	1	Black/African American
394974	356303	1	Black/African American

[394975 rows x 3 columns]

[4]: data_stops_raw

```
[4]:
```

	stop_id	ori	agency	exp_years	date_stop	time_stop	\
0	2443	CA0371100	SD	10	2018-07-01	00:01:37	
1	2444	CA0371100	SD	18	2018-07-01	00:03:34	
2	2447	CA0371100	SD	1	2018-07-01	00:05:43	
3	2447	CA0371100	SD	1	2018-07-01	00:05:43	
4	2448	CA0371100	SD	3	2018-07-01	00:19:06	
...	
391129	356019	CA0371100	SD	1	2020-09-30	23:05:00	
391130	356025	CA0371100	SD	1	2020-09-30	23:38:00	
391131	356080	CA0371100	SD	18	2020-09-30	15:31:00	
391132	356300	CA0371100	SD	18	2020-09-30	19:30:00	
391133	356303	CA0371100	SD	1	2020-09-30	19:37:52	

	stopduration	stop_in_response_to_cfs	officer_assignment_key	\
0	30	0	1	
1	10	0	1	
2	15	1	10	
3	15	1	10	
4	5	0	1	
...	
391129	7	1	1	
391130	30	1	1	
391131	5	0	1	
391132	180	1	1	
391133	45	0	1	

	assignment	...	\
0	Patrol, traffic enforcement, field operations	...	
1	Patrol, traffic enforcement, field operations	...	
2	Other	...	
3	Other	...	
4	Patrol, traffic enforcement, field operations	...	
...	
391129	Patrol, traffic enforcement, field operations	...	
391130	Patrol, traffic enforcement, field operations	...	
391131	Patrol, traffic enforcement, field operations	...	
391132	Patrol, traffic enforcement, field operations	...	
391133	Patrol, traffic enforcement, field operations	...	

	beat_name	pid	isstudent	perceived_limited_english	\
0	Pacific Beach 122	1	0	0	
1	Mission Beach 121	1	0	0	

2	El Cerrito	822	1	0	0
3	El Cerrito	822	2	0	0
4	Ocean Beach	614	1	0	0
...
391129	Harborview	527	1	0	0
391130	Core-Columbia	524	1	0	0
391131	Unknown	999	1	0	0
391132	Carmel Mountain	232	1	0	0
391133	Golden Hill	517	1	0	0

	perceived_age	perceived_gender	gender_nonconforming	gend	gend_nc	\
0	25	Male		0 1	NaN	
1	25	Male		0 1	NaN	
2	30	Male		0 1	NaN	
3	30	Female		0 2	NaN	
4	23	Male		0 1	NaN	
...
391129	50	Female		0 2	NaN	
391130	35	Male		0 1	NaN	
391131	60	Male		0 1	NaN	
391132	25	Male		0 1	NaN	
391133	28	Male		0 1	NaN	

	perceived_lgbt
0	No
1	No
2	No
3	No
4	No
...	...
391129	No
391130	No
391131	No
391132	No
391133	No

[391134 rows x 29 columns]

```
[5]: data_census_race
```

```
[5]:
percentage of population
White 45.0
Black/African American 5.5
Hispanic/Latino/a 34.1
Middle Eastern or South Asian 0.9
Asian 12.6
Native American 1.3
```

Data Wrangling: How do we make our data useable

As of now, our raw data has a lot of information that we don't really need to answer our data science question. (For the SDPD in 2019, is there is significant difference in the likely hood of someone being stopped according to their race?) The columns that we need consists of the **rac**es of the people stopped and the **date**. (to restrict our time interval) In order to do this and combine it into one table, we need to first drop all the columns in the stops data we aren't using and merge that table to the race table at the **stop_id**.

Then we have to remove duplicates that arise from the merge. The reason why the merge creates duplicates is because multiple people can be stopped at one stop_id represented by the **pid**(person id) which means that there will be an addition copy of the stop_id from the dates table. (Has unique stops and doesn't account for pid) In addition to removing duplicates, we have to set the final table to only contain stop_ids with a date in 2019. To this, we need to convert the dates in the date column to something we can read, such as a **datetime**. After converting the date column of strings to a datetime, we can restrict the table to only include stops from the year 2019.

```
[6]: data_race = data_race_raw.set_index('stop_id')
data_race
```

```
[6]:
```

	pid	race
stop_id		
2443	1	White
2444	1	White
2447	1	Hispanic/Latino/a
2447	2	Hispanic/Latino/a
2448	1	White
...
356019	1	Black/African American
356025	1	Black/African American
356080	1	White
356300	1	Black/African American
356303	1	Black/African American

[394975 rows x 2 columns]

```
[7]: data_date = pd.DataFrame().assign(date = data_stops_raw.get('date_stop'),
    ↳ stop_id = data_stops_raw.get('stop_id')).set_index('stop_id')
data_date
```

```
[7]:
```

	date
stop_id	
2443	2018-07-01
2444	2018-07-01
2447	2018-07-01
2447	2018-07-01

```

2448      2018-07-01
...
356019    2020-09-30
356025    2020-09-30
356080    2020-09-30
356300    2020-09-30
356303    2020-09-30

```

[391134 rows x 1 columns]

```

[8]: #Merge race data set with the dates from the stop data set with the stop_id
data_merged = data_race.merge(data_date, left_index = True, right_index = True)
data_merged

```

```

[8]:      pid      race      date
stop_id
2443      1      White  2018-07-01
2444      1      White  2018-07-01
2447      1  Hispanic/Latino/a  2018-07-01
2447      1  Hispanic/Latino/a  2018-07-01
2447      2  Hispanic/Latino/a  2018-07-01
...
356019      1  Black/African American  2020-09-30
356025      1  Black/African American  2020-09-30
356080      1      White  2020-09-30
356300      1  Black/African American  2020-09-30
356303      1  Black/African American  2020-09-30

```

[595128 rows x 3 columns]

```

[9]: #Remove Duplicates
data_final = data_merged.drop_duplicates()

#Get the year from the date string
def to_year(date):
    dt = datetime.datetime.strptime(date, '%Y-%m-%d')
    return dt.year

#Include data within subjected time interval
data_final = data_final[data_final.get('date').apply(to_year) == 2019]
data_final

```

```

[9]:      pid      race      date
stop_id
84362      1  Hispanic/Latino/a  2019-01-01
84364      1      White  2019-01-01
84369      1  Black/African American  2019-01-01

```

84372	2	Hispanic/Latino/a	2019-01-01
84376	1	Middle Eastern or South Asian	2019-01-01
...
254761	8	White	2019-12-31
254771	2	White	2019-12-31
254776	1	Native American	2019-12-31
255002	4	White	2019-12-31
255002	5	White	2019-12-31

[8398 rows x 3 columns]

Exploratory/Descriptive Analysis: What does our data say?

Now that we have our data neated sorted and ready to use, lets use that to answer our question. In order to do that we started off by producing some basic descriptive analysis to get a general idea of the trend of our data. For the descriptive analysis, we decided to compare the percentages of each race stopped versus the percentages of each race's demographic to see if there was any significant difference between the two. To do this, we created a new table showing the percentage of each race stopped and then merged the demographic table to that table. Then generate basic statistics that comparet the two distributions such as the absolute mean difference, the standard deviation of the absolute differences, and the range of the absolute differences.

```
[10]: #Generate Race Percentages Table from the data wrangled data_final
race_percentage = data_final.groupby('race').count()/data_final.shape[0]*100
race_percentage = race_percentage.drop(columns = ['date']).
    ↳rename(columns={'pid':'percentage stopped'})
race_percentage
```

```
[10]:
percentage stopped
race
Asian                10.431055
Black/African American  20.159562
Hispanic/Latino/a      22.552989
Middle Eastern or South Asian  8.001905
Native American        3.346035
Pacific Islander        6.215766
White                 29.292689
```

Now we see that there's a difference between the two distributions in the following table below.

```
[11]: #Now merge the census data and sort by lowest population to highest
race_census_percentage = race_percentage.merge(data_census_race,left_index =_
    ↳True,right_index = True)
race_census_percentage = race_census_percentage.sort_values('percentage of_
    ↳population', ascending = True)
race_census_percentage
```

	percentage stopped	percentage of population
Pacific Islander	6.215766	0.6
Middle Eastern or South Asian	8.001905	0.9
Native American	3.346035	1.3
Black/African American	20.159562	5.5
Asian	10.431055	12.6
Hispanic/Latino/a	22.552989	34.1
White	29.292689	45.0

Basic Statistics generated below:

```
[12]: difference = race_census_percentage.get('percentage_
↳stopped')-race_census_percentage.get('percentage of population')
mean_abs_diff = abs(difference).mean()
std_abs_diff = np.std(abs(difference))
range_abs_diff = abs(difference).max()-abs(difference).min()
print('mean abs diff: ' + str(mean_abs_diff) + '%, std abs diff: ' +
↳str(std_abs_diff) + '%, range abs diff: ' + str(range_abs_diff)+'%')
```

```
mean abs diff: 8.406647841322764%, std abs diff: 5.22518917771041%, range abs
diff: 13.661276494403433%
```


Our descriptive analysis shows that there's a decent chance that race is correlated with the likelihood of being stopped by the police in San Diego with the mean difference from the demographic being 8.41%. Assuming that these results are reliable(proven later in our statistical analysis), we can then use this to justify or disprove our hypothesis that blacks and hispanics are stopped more frequently by the police compared to other races.

Data Visualization: What does the data show us

Once we completed some basic analysis, the next step is to visualize the data and see how each races compares to it's demographic. With Pandas it's really easy and simple to perform since we already have a completed table from our exploratory analysis and data wrangling we did. We decided to make a grouped bar graph to compare the distribution of two categorical variables. We decided to choose blue and orange as the colors and make the graph a bit wider to make it easier on the eyes. As you can see, minorities are stopped much more often than the rest of the races, especially in Black/African American and much less in Asian and White which supports our hypothesis. However, Hispanics/Latinos are surprisingly stopped less relative to their population according to the visualization which goes against our hypothesis.

```
[13]: race_census_percentage.plot(kind = 'bar', figsize = [8,5],color =
↳['orange','royalblue'])
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2c927fe358>
```

Statistical Analysis: How do we know that our results are trustworthy

After performing exploratory analysis to generate basic statistics and producing a visualization, we need to make sure that these results are valid and significant. In order to achieve this we will perform a hypothesis test and calculate a confidence interval and p-value for our statistic.

- Null: There is no significant difference between the percentage of races stopped respective to their demographic
- Alternate: There is a significant difference between the percentage of race stopped respective to their demographic

[14]: *#Test Statistic will be the Mean Difference*

```
test_stat = mean_abs_diff
test_stat
```

[14]: 8.406647841322764

[15]: *#We'll generate about 5000 sample test stats using the census data to create a 95% confidence interval*

```

num_repetitions = 5000
population = data_final.shape[0]

simulated_test_stats = np.array([])

for i in range(num_repetitions):
    model_proportions = race_census_percentage.get('percentage of population')/
    ↪100
    sample = np.random.multinomial(population, model_proportions)/population
    sim_test_stat = abs(model_proportions-sample).mean()*100
    simulated_test_stats = np.append(simulated_test_stats, sim_test_stat)

simulated_test_stats

```

```

[15]: array([0.27462321, 0.16178002, 0.26500868, ..., 0.1989045 , 0.1827646 ,
           0.25130473])

```

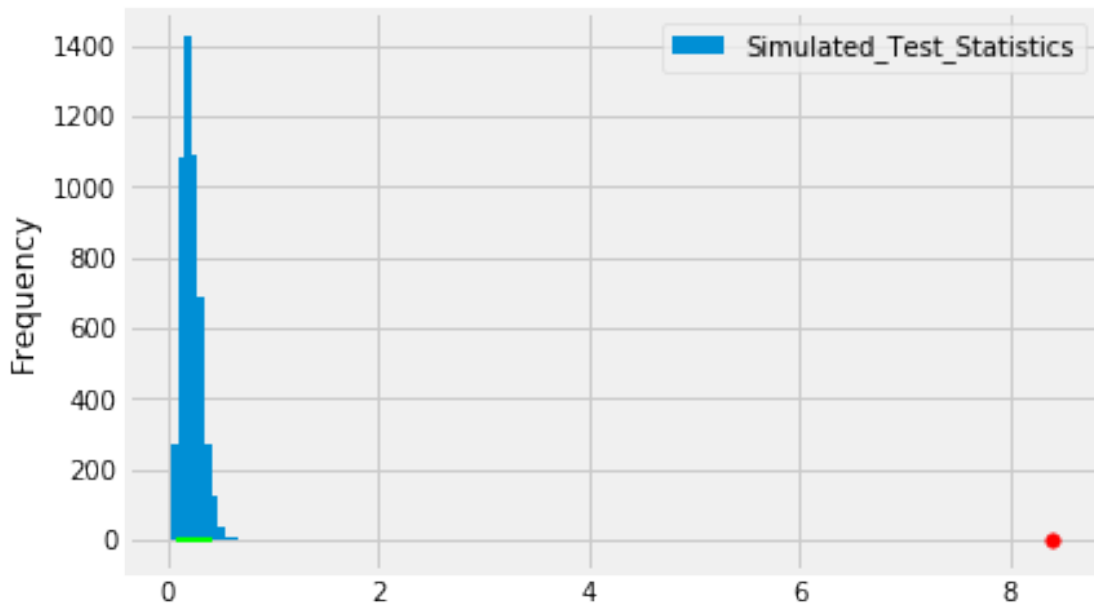
```

[16]: #Lets look at the distribution and generate the 95% confidence interval
t = pd.DataFrame().assign(Simulated_Test_Statistics = simulated_test_stats)
t.plot(kind='hist')

confidence_interval = [np.percentile(simulated_test_stats,2.5),np.
    ↪percentile(simulated_test_stats,97.5)]
plt.scatter(test_stat, 0, color='red', s=30);
plt.plot(confidence_interval,[0,0], color = 'lime', linewidth = 2)
print('Confidence Interval: [' + str(confidence_interval[0]) + ', ' +
    ↪str(confidence_interval[1]) + ']')

```

Confidence Interval: [0.08422379478106999, 0.4320547409247092]



```
[17]: #Now lets generate a p value
p_value = np.count_nonzero(simulated_test_stats >= test_stat)/
    ↳simulated_test_stats.shape[0]
print('p value: ' + str(p_value))
```

p value: 0.0

We reject the null, therefore the difference in the percentage of races being stopped is statistically significant. With this test, we are confident to use the results as justification to prove if our hypothesis is correct or incorrect.

Geospatial Analysis: Is there any bias in our data due to the location of the stops

In this analysis we compare the total distribution of the San Diego police stops to the total demographic of San Diego to see if there's any difference between the two. However, the police may not uniformly patrol each district of San Diego. For example, the police might be more active in urban areas where minorities may make up more of the population as opposed to suburbs where most of the populations are whites, hispanics, and asians. We'll perform a geospatial analysis to see if the frequency a police stops at a location has an effect on the mean difference of races to demographic stopped. This will tell us if there's any bias in our data and how severe it is.

The process will involve using geopandas to read a geodataframe in the form of a shp file. The shape file data we will be using is from [ripa](#) which contains a beat map of the san diego county. (A beat is a territory an officer patrols) To start off with our geospatial analysis, we need to do a bit of addition data wrangling by adding the beat data to our data final. After adding all the necessary columns (The frequency of stops at each beat and the mean absolute difference of race to demographic of each beat), we'll use the beat map to create visualizations to look for possible correlations, hence bias in our data. More specifically, we choose to make a cartogram that represented how much each beat was affected by the police according to the frequency of stops and

a choropleth that compared the mean absolute difference of each beat.

```
[18]: #Data Wrangling

#We'll add the location to our final data set
data_loc = pd.DataFrame().assign(stop_id = data_stops_raw.get('stop_id'), beat =
↳data_stops_raw.get('beat'))
data_loc = data_loc.set_index('stop_id')
data_geo = data_final.merge(data_loc, left_index = True, right_index = True)
data_geo = data_geo.drop_duplicates()
data_geo
```

```
[18]:
```

	pid	race	date	beat
stop_id				
84362	1	Hispanic/Latino/a	2019-01-01	839
84364	1	White	2019-01-01	124
84369	1	Black/African American	2019-01-01	614
84372	2	Hispanic/Latino/a	2019-01-01	122
84376	1	Middle Eastern or South Asian	2019-01-01	122
...
254761	8	White	2019-12-31	521
254771	2	White	2019-12-31	512
254776	1	Native American	2019-12-31	121
255002	4	White	2019-12-31	614
255002	5	White	2019-12-31	614

[8398 rows x 4 columns]

```
[19]: #This is the frequency of stops for each beat
geo_final = data_geo.groupby('beat').count()
geo_final = geo_final.assign(frequency = geo_final.get('pid')).
↳drop(columns=['pid', 'date', 'race'])
geo_final
```

```
[19]:
```

	frequency
beat	
111	67
112	26
113	23
114	50
115	53
...	...
934	61
935	19
936	10
937	20
999	69

[122 rows x 1 columns]

```
[20]: #This is the returns absolute mean difference of each beat
def get_mean_diff(races):
    race_percentages = races/races.sum()*100
    return abs(race_percentages.get('frequency')- data_census_race.
    ↳get('percentage of population')).mean()
```

```
[22]: #Generating the final table with mean abs difference and frequency
geo_diff = data_geo.groupby(['beat', 'race']).count().unstack(fill_value=0).
    ↳stack()
geo_diff = geo_diff.assign(frequency = geo_diff.get('pid')).
    ↳drop(columns=['pid', 'date'])

mean_diff = np.array([])
for i in geo_final.index:
    mean_diff = np.append(mean_diff, get_mean_diff(geo_diff.loc[i]))

geo_final = geo_final.assign(mean_diff = mean_diff)
geo_final
```

```
[22]:      frequency  mean_diff
beat
111          67  10.233262
112          26   9.791209
113          23   7.959006
114          50  11.171429
115          53  11.772507
...
934          61  10.422014
935          19   7.278195
936          10   7.571429
937          20  14.200000
999          69  10.591718
```

[122 rows x 2 columns]

```
[33]: beat_map_raw = gpd.read_file('pd_beats_datasd/pd_beats_datasd.shp')
beat_map_raw
```

```
[33]:      objectid  beat  div  serv      name \
0           3   935   9   930  NORTH CITY
1           7     0   0     0   SAN DIEGO
2           8   511   5   510      None
3           9   722   7   720   NESTOR
4          10   314   3   310  BIRDLAND
```

```

..      ...      ...      ...      ...
135      610      243      2      240      MIRAMAR
136      616      937      9      930      BLACK MOUNTAIN RANCH
137      617      936      9      930      TORREY HIGHLANDS
138      618      233      2      230      RANCHO PENASQUITOS
139      619      235      2      230      SAN PASQUAL

```

```

                                geometry
0      MULTIPOLYGON (((6268975.465 1931147.469, 62689...
1      MULTIPOLYGON (((6261648.576 1836846.672, 62616...
2      MULTIPOLYGON (((6261640.429 1836823.561, 62616...
3      POLYGON ((6302781.000 1793246.001, 6302905.000...
4      POLYGON ((6284667.652 1874418.895, 6284694.392...
..
135     POLYGON ((6295777.084 1908435.558, 6295790.126...
136     POLYGON ((6295209.132 1952772.422, 6295208.647...
137     POLYGON ((6287497.113 1936631.116, 6287462.453...
138     POLYGON ((6306176.245 1943421.146, 6306158.416...
139     POLYGON ((6338763.812 1971222.216, 6338683.658...

```

[140 rows x 6 columns]

```

[55]: beat_visual= beat_map_raw.merge(geo_final, left_on = 'beat',right_index = True,
    ↪how = 'left')
beat_visual

```

```

[55]:      objectid  beat  div  serv      name \
0           3    935    9    930      NORTH CITY
1           7     0    0     0      SAN DIEGO
2           8    511    5    510           None
3           9    722    7    720      NESTOR
4          10    314    3    310      BIRDLAND
..
135         610    243    2    240      MIRAMAR
136         616    937    9    930  BLACK MOUNTAIN RANCH
137         617    936    9    930    TORREY HIGHLANDS
138         618    233    2    230    RANCHO PENASQUITOS
139         619    235    2    230    SAN PASQUAL

```

```

                                geometry  frequency  mean_diff
0      MULTIPOLYGON (((6268975.465 1931147.469, 62689...      19.0      7.278195
1      MULTIPOLYGON (((6261648.576 1836846.672, 62616...      NaN      NaN
2      MULTIPOLYGON (((6261640.429 1836823.561, 62616...     117.0     11.084737
3      POLYGON ((6302781.000 1793246.001, 6302905.000...      27.0     13.114286
4      POLYGON ((6284667.652 1874418.895, 6284694.392...      21.0     20.136054
..
135     POLYGON ((6295777.084 1908435.558, 6295790.126...      23.0     12.662112

```

```

136 POLYGON ((6295209.132 1952772.422, 6295208.647...      20.0  14.200000
137 POLYGON ((6287497.113 1936631.116, 6287462.453...      10.0   7.571429
138 POLYGON ((6306176.245 1943421.146, 6306158.416...      39.0  13.915018
139 POLYGON ((6338763.812 1971222.216, 6338683.658...       2.0  15.714286

```

[140 rows x 8 columns]

Cartogram: As you can see below, the downtown areas have a lot more police activity

```

[35]: #ax = geoplot.cartogram(
#     beat_vis, scale='frequency', limits=(0.2, 1),
#     edgecolor='None', figsize=(7, 8)
#)
#geoplot.polyplot(beat_vis, edgecolor='gray', ax=ax

```

Choropleth: As you can see below, beats from the downtown area of San Diego County have higher mean differences compared to the suburbs more north.

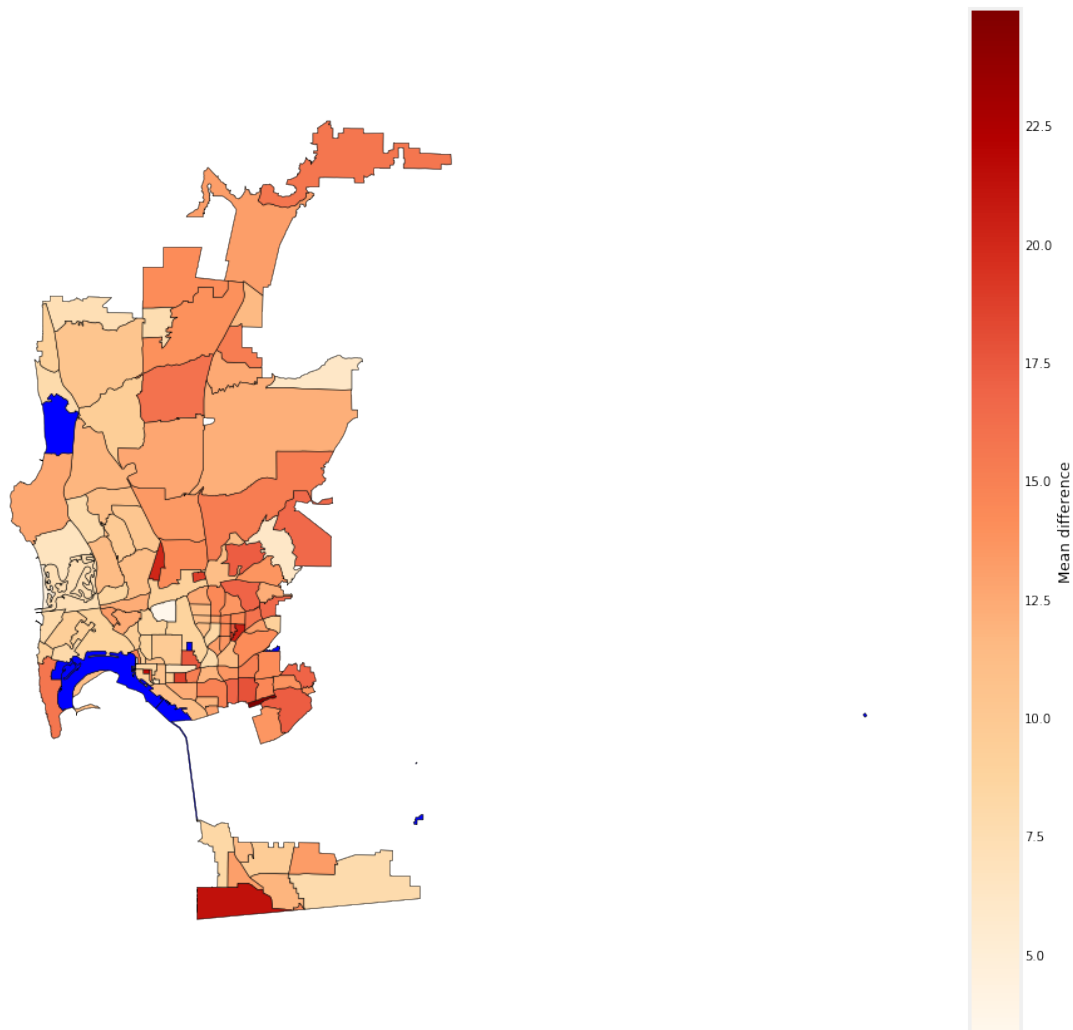
```

[82]: beat_visual.plot(column = 'mean_diff',figsize = [15,15], legend=True,
    cmap='OrRd', edgecolor="black",
    legend_kwds={'label': "Mean difference"},
    missing_kwds={'color': 'blue'})
plt.axis('off')
print('Blue Represents that the beat has no stops')

plt.suptitle('San Diego 2019: Mean Difference of races stopped vs. Demographic')
plt.show()

```

Blue Represents that the beat has no stops



Overlay Visualization: Now lets see how similar the two graphs compare when we overlay them

[]:

Conclusion: Is our question answered?

Although we performed a hypothesis test and recieved a p-value of less than .05, our exploratory analysis may still contain biases and inaccuracies as shown in our geospatial analysis. However we believed that this isn't enough to completely invalidate our findings, but be conscious that our results maybe slightly skewed from not accounting for societal bias.(Police are more active in downtown areas which are typically inhabited by more minioritise) From our hypothesis test in our statistical analysis, we conclude that there's a relationship between someone's race and their likelihood of being stopped by the people. Parts of our hypothesis were proven right and wrong.

We said that blacks/african americans were more likely of being stopped, which the data firmly suggested. However we also hypothesized that latinos/hispanics would also be targeted more, but according to our findings, this wasn't the case as it was even less likely to be stopped if one was a hispanic. In conclusion the trends to seems to be that minorities are stopped more often than other races in the demographic of San Diego in 2019.

```
[83]: print('Congrats for making it to the END! :)')
```

```
Congrats for making it to the END! :)
```